

Say, Echo, Do

Reading the gap between what markets hear and what money does

A concept explainer · [Author Name] · September 2026

Summary. Markets hear three voices. Institutions *say* things. The media *echoes* them. Capital *does* things. Most sentiment models listen only to the first two. This explainer describes a framework that listens for *disagreement* between the voices, using AI embeddings to read what the narrative really means and quantitative finance to read what money is really doing. It rests on two inversions: *trust deeds over words*, and *fade the echo, follow the news*.

1 The Idea

Everyone in markets has seen this pattern. A respected bank publishes a gloomy call. Headlines multiply. The stock drops, but less than the noise suggests, and weeks later it is higher than before. Looking back, someone was buying all along.

It is tempting to call this a conspiracy. We don't need to. We only need to *measure the disagreement* between three voices:

Voice	What it is	Where we see it
Say	What institutions publicly claim	Ratings, target changes, research notes, management tone, interviews
Echo	What the media repeats	News volume, syndication, social posts, emotional intensity
Do	What capital actually does	Holdings, futures positioning, insider trades, options flow, price




Figure 1: The false-alarm pattern (stylised). Positioning builds (1) before the alarm (2); the price falls far less than the alarm implies (3), then recovers (4).

The two inversions

- 1. Deeds over words.** When **Say** and **Do** disagree, trust **Do**. A bearish institutional call accompanied by institutional buying is a bullish signal.
- 2. Fade the echo, follow the news.** Split media sentiment into *new information* and *echo*. Prices tend to drift after real news; it is the amplification that overshoots.

The price is the referee. If the story says the stock should fall 6% and it falls 2%, someone absorbed the selling. And the logic runs both ways: euphoric coverage while institutions quietly sell is the mirror image at market tops. The framework never assumes “always be contrarian”; it asks which of four situations we are in.

	Narrative is bearish	Narrative is bullish
Money & price oppose the story	False alarm Fade the fear: buy	False euphoria Fade the hype: sell
Money & price confirm the story	Real bad news Follow it: sell	Real good news Follow it: buy

Figure 2: The decision matrix. The narrative alone never decides the trade; its agreement with money and price does.

2 Why Inverting Sentiment Can Work

The ingredients are well established. Media pessimism creates short-term pressure that later reverses [12]. Attention-grabbing news pulls in less-informed buyers [2]. Sentiment can push prices away from value because betting against it is risky [7, 11]. Sophisticated traders profit by trading against forced or emotional flow [4]. Yet prices also *drift* after genuine news [5], which is why a naive contrarian rule fails.

A tiny model makes the logic precise. A stock is worth v ; a narrative $m = v + \text{noise}$ arrives; the crowd moves the price by φ per unit of narrative, while a rational observer would use weight $w^* = \text{Var}(v) / \text{Var}(m)$.




Figure 3: Result 1. Whether to follow or fade the narrative depends on one comparison: how strongly the crowd reacts versus how much the narrative is actually worth.

The model gives three results:

Result 1. Fade the story *only* when the crowd overreacts (Figure 3).

Result 2. The price move *net of* what the narrative implied reveals the hidden informed trader, and it predicts continuation in its own direction.

Result 3. Any observation of real positioning, however noisy, makes the forecast sharper.

3 Embeddings: A Map of Meaning

An embedding model turns text into a point in a high-dimensional space, so that texts with similar meaning land close together. Think of it as a *map of meaning*. On this map, **directions carry concepts** (moving one way makes text more alarming) and **distances carry surprise** (how far the story landed from where it was expected to land).




Figure 4: The map of meaning (a 2D cartoon of a 1024-dimensional space). Stories cluster by topic. Sentiment is a direction from calm to alarm. Surprise is the gap between the expected and the actual story, and *sentiment surprise* is the part of that gap along the alarm direction.

3.1 Teaching the map what matters to markets

Off-the-shelf embeddings organise text by *topic*: two oil stories sit together whether they turned out bullish or bearish. We fine-tune the model on history so that stories with similar *market aftermaths* move together (Figure 5). After training, “this kind of news tends to be overdone” becomes a place on the map.

The math: scoring a story

Each document becomes an embedding z . A small sequence model predicts the *expected* next story $\hat{z}$ from recent coverage and market conditions. Then

$$\underbrace{\varepsilon = z - \hat{z}}_{\text{surprise}}, \quad \underbrace{s = \langle w_s, \varepsilon \rangle}_{\text{sentiment surprise}}, \quad w_s \text{ with topic directions removed.}$$

Fine-tuning uses a contrastive loss that pulls together stories whose subsequent returns, volume and volatility were similar [10].




Figure 5: What fine-tuning does. Before, stories cluster by subject and good and bad outcomes are mixed. After, stories cluster by what happened next, across subjects.

In plain words. We score news on how much worse or better it is than what the market was bracing for, not on how negative it sounds.

4 Inversion 1: Fade the Echo, Follow the News

When one story appears and forty copies follow, the copies add attention but no information. We model story arrivals as a *self-exciting* process: each story briefly raises the chance of more stories [1, 9]. The model then tells us, for every story, the probability that it was *new* rather than *triggered* by earlier coverage.



Figure 6: Echo versus news. Top: each new story (blue) can trigger a cascade of echoes (red). Bottom: the story rate jumps with every arrival and decays back to the baseline. The first episode is mostly echo; the later ones are mostly news.

The math: splitting sentiment

Story arrivals follow $\lambda(t) = \mu + \sum_{\tau_j < t} \alpha e^{-\beta(t-\tau_j)}$. The probability that story j was new is $p_j = \mu/\lambda(\tau_j)$, which splits total sentiment into

$$s^{\text{new}} = \sum_j p_j s_j \quad (\text{follow}), \quad s^{\text{echo}} = \sum_j (1 - p_j) s_j \quad (\text{fade}).$$

The ratio α/β is the share of coverage that is echo; near 1, the narrative is feeding on itself.

5 Inversion 2: Deeds over Words

The most useful comparisons match the *same kind of actor* on both sides:

Actor	Say	Do
Management	Tone on the earnings call	Insider purchases and sales
Asset managers	Public outlooks and interviews	Quarterly holdings disclosures
Futures participants	Published macro views	Weekly positioning reports
The market	Analyst ratings and targets	Options flow, skew, fund flows, short interest

Plot every event by what was said against what was done (Figure 7). Events near the diagonal are consistent: follow them. Events far off the diagonal are where the signal lives.




Figure 7: The Say–Do plane. The distance from the diagonal is the Say–Do gap. Top-left is the false-alarm zone; bottom-right is the false-euphoria zone.

Two technical pieces make this work:

- **Recovering hidden positioning.** Disclosures arrive at different speeds: insider trades within days, futures positions weekly, fund holdings quarterly and weeks late. A *Kalman filter* treats true positioning as a hidden quantity and each disclosure as a delayed, noisy snapshot of it. It outputs the best current estimate *and* how uncertain it is, so stale data is automatically trusted less.
- **A credibility ledger.** Every source keeps a running track record. Reliable voices get more weight; reliably *wrong* voices are not ignored, their sign is flipped. The inversion is earned from data, not assumed.

6 Reading the Price

6.1 Absorption: how much the price refused to move

The map of meaning lets us ask: *when similar stories hit in similar conditions, how did prices react?* Averaging the reactions of the nearest historical analogues gives the *expected* reaction. Absorption is the gap between expected and actual (Figure 8).




Figure 8: Absorption. The story implied a 6% fall; the stock fell 1.9%. The missing 4.1% is the footprint of a buyer.

A second, independent check comes from market microstructure. Trading a quantity Q typically moves the price by about $\sigma\sqrt{Q/V}$, the well-documented *square-root law* [3]. Heavy panic selling with little price damage means high effort and a small result; someone absorbed the flow.

6.2 Who moved first?

The strongest false alarms have a clear order: money moves, *then* the story turns against it. Plot positioning against headline gloom over the past month and trace the path. If the loop turns one way, money led; the other way, money followed. The *signed area* of the loop, a tool from rough-path theory [6], measures how strongly.




Figure 9: Lead-lag as a loop. The direction of travel around the loop tells us who moved first; the enclosed area tells us how strongly.

7 Putting It Together

All signals feed a regime model that decides, event by event, which world we are in: *informational* (money confirms the story: follow), *overreaction* (mostly echo, no opposition: fade gently), or *strategic* (money opposes the story and moved first: fade firmly). Each forecast comes with a calibrated confidence range [8]; a wide range means a small position.




Figure 10: The whole framework. Three voices become three families of signals, which a regime model turns into a single decision.

Illustration: a hypothetical downgrade

Before. Over eight days, call options build up, put skew flattens, and two directors buy shares.

Day 0. Two brokers downgrade the stock; coverage jumps from 15 stories a day to 140.

Signal	Reading
Sentiment surprise (new / echo)	-0.5σ / -2.3σ : mostly echo
Say-Do gap	bearish words, bullish positioning
Expected vs actual reaction	-6.0% vs -1.9% : strong absorption
Lead-lag loop	positive: money moved first
Regime (informational / overreaction / strategic)	10% / 28% / 62%

Reading. A probable false alarm. The framework leans long, but sizes modestly because the confidence range still includes losses.

8 What This Is, and What It Isn't

It measures disagreement, not intent. A “strategic” reading means money and narrative moved in opposite directions in a particular order. It is not evidence that anyone lied, and large firms keep information barriers between research and trading.

It is not blind contrarianism. Real news is followed; only echo and contradicted narratives are faded.

It uses only public information. Creating or spreading narratives to move prices would itself be market manipulation. The framework reads the conversation; it never joins it.

It is a concept, not yet a proven strategy. Every component must be tested on point-in-time data before any capital is involved.

References

- [1] Bacry, E., Mastromatteo, I., and Muzy, J.-F. (2015). Hawkes processes in finance. *Market Microstructure and Liquidity*, 1(1).
- [2] Barber, B. M., and Odean, T. (2008). All that glitters: The effect of attention and news on the buying behavior of individual and institutional investors. *Review of Financial Studies*, 21(2), 785–818.
- [3] Bouchaud, J.-P., Bonart, J., Donier, J., and Gould, M. (2018). *Trades, Quotes and Prices*. Cambridge University Press.
- [4] Brunnermeier, M. K., and Pedersen, L. H. (2005). Predatory trading. *Journal of Finance*, 60(4), 1825–1863.
- [5] Chan, W. S. (2003). Stock price reaction to news and no-news: Drift and reversal after headlines. *Journal of Financial Economics*, 70(2), 223–260.
- [6] Chevyrev, I., and Kormilitzin, A. (2016). A primer on the signature method in machine learning. arXiv:1603.03788.
- [7] De Long, J. B., Shleifer, A., Summers, L. H., and Waldmann, R. J. (1990). Noise trader risk in financial markets. *Journal of Political Economy*, 98(4), 703–738.
- [8] Gibbs, I., and Candès, E. (2021). Adaptive conformal inference under distribution shift. *NeurIPS*, 34.
- [9] Hawkes, A. G. (1971). Spectra of some self-exciting and mutually exciting point processes. *Biometrika*, 58(1), 83–90.
- [10] van den Oord, A., Li, Y., and Vinyals, O. (2018). Representation learning with contrastive predictive coding. arXiv:1807.03748.
- [11] Shleifer, A., and Vishny, R. W. (1997). The limits of arbitrage. *Journal of Finance*, 52(1), 35–55.
- [12] Tetlock, P. C. (2007). Giving content to investor sentiment: The role of media in the stock market. *Journal of Finance*, 62(3), 1139–1168.